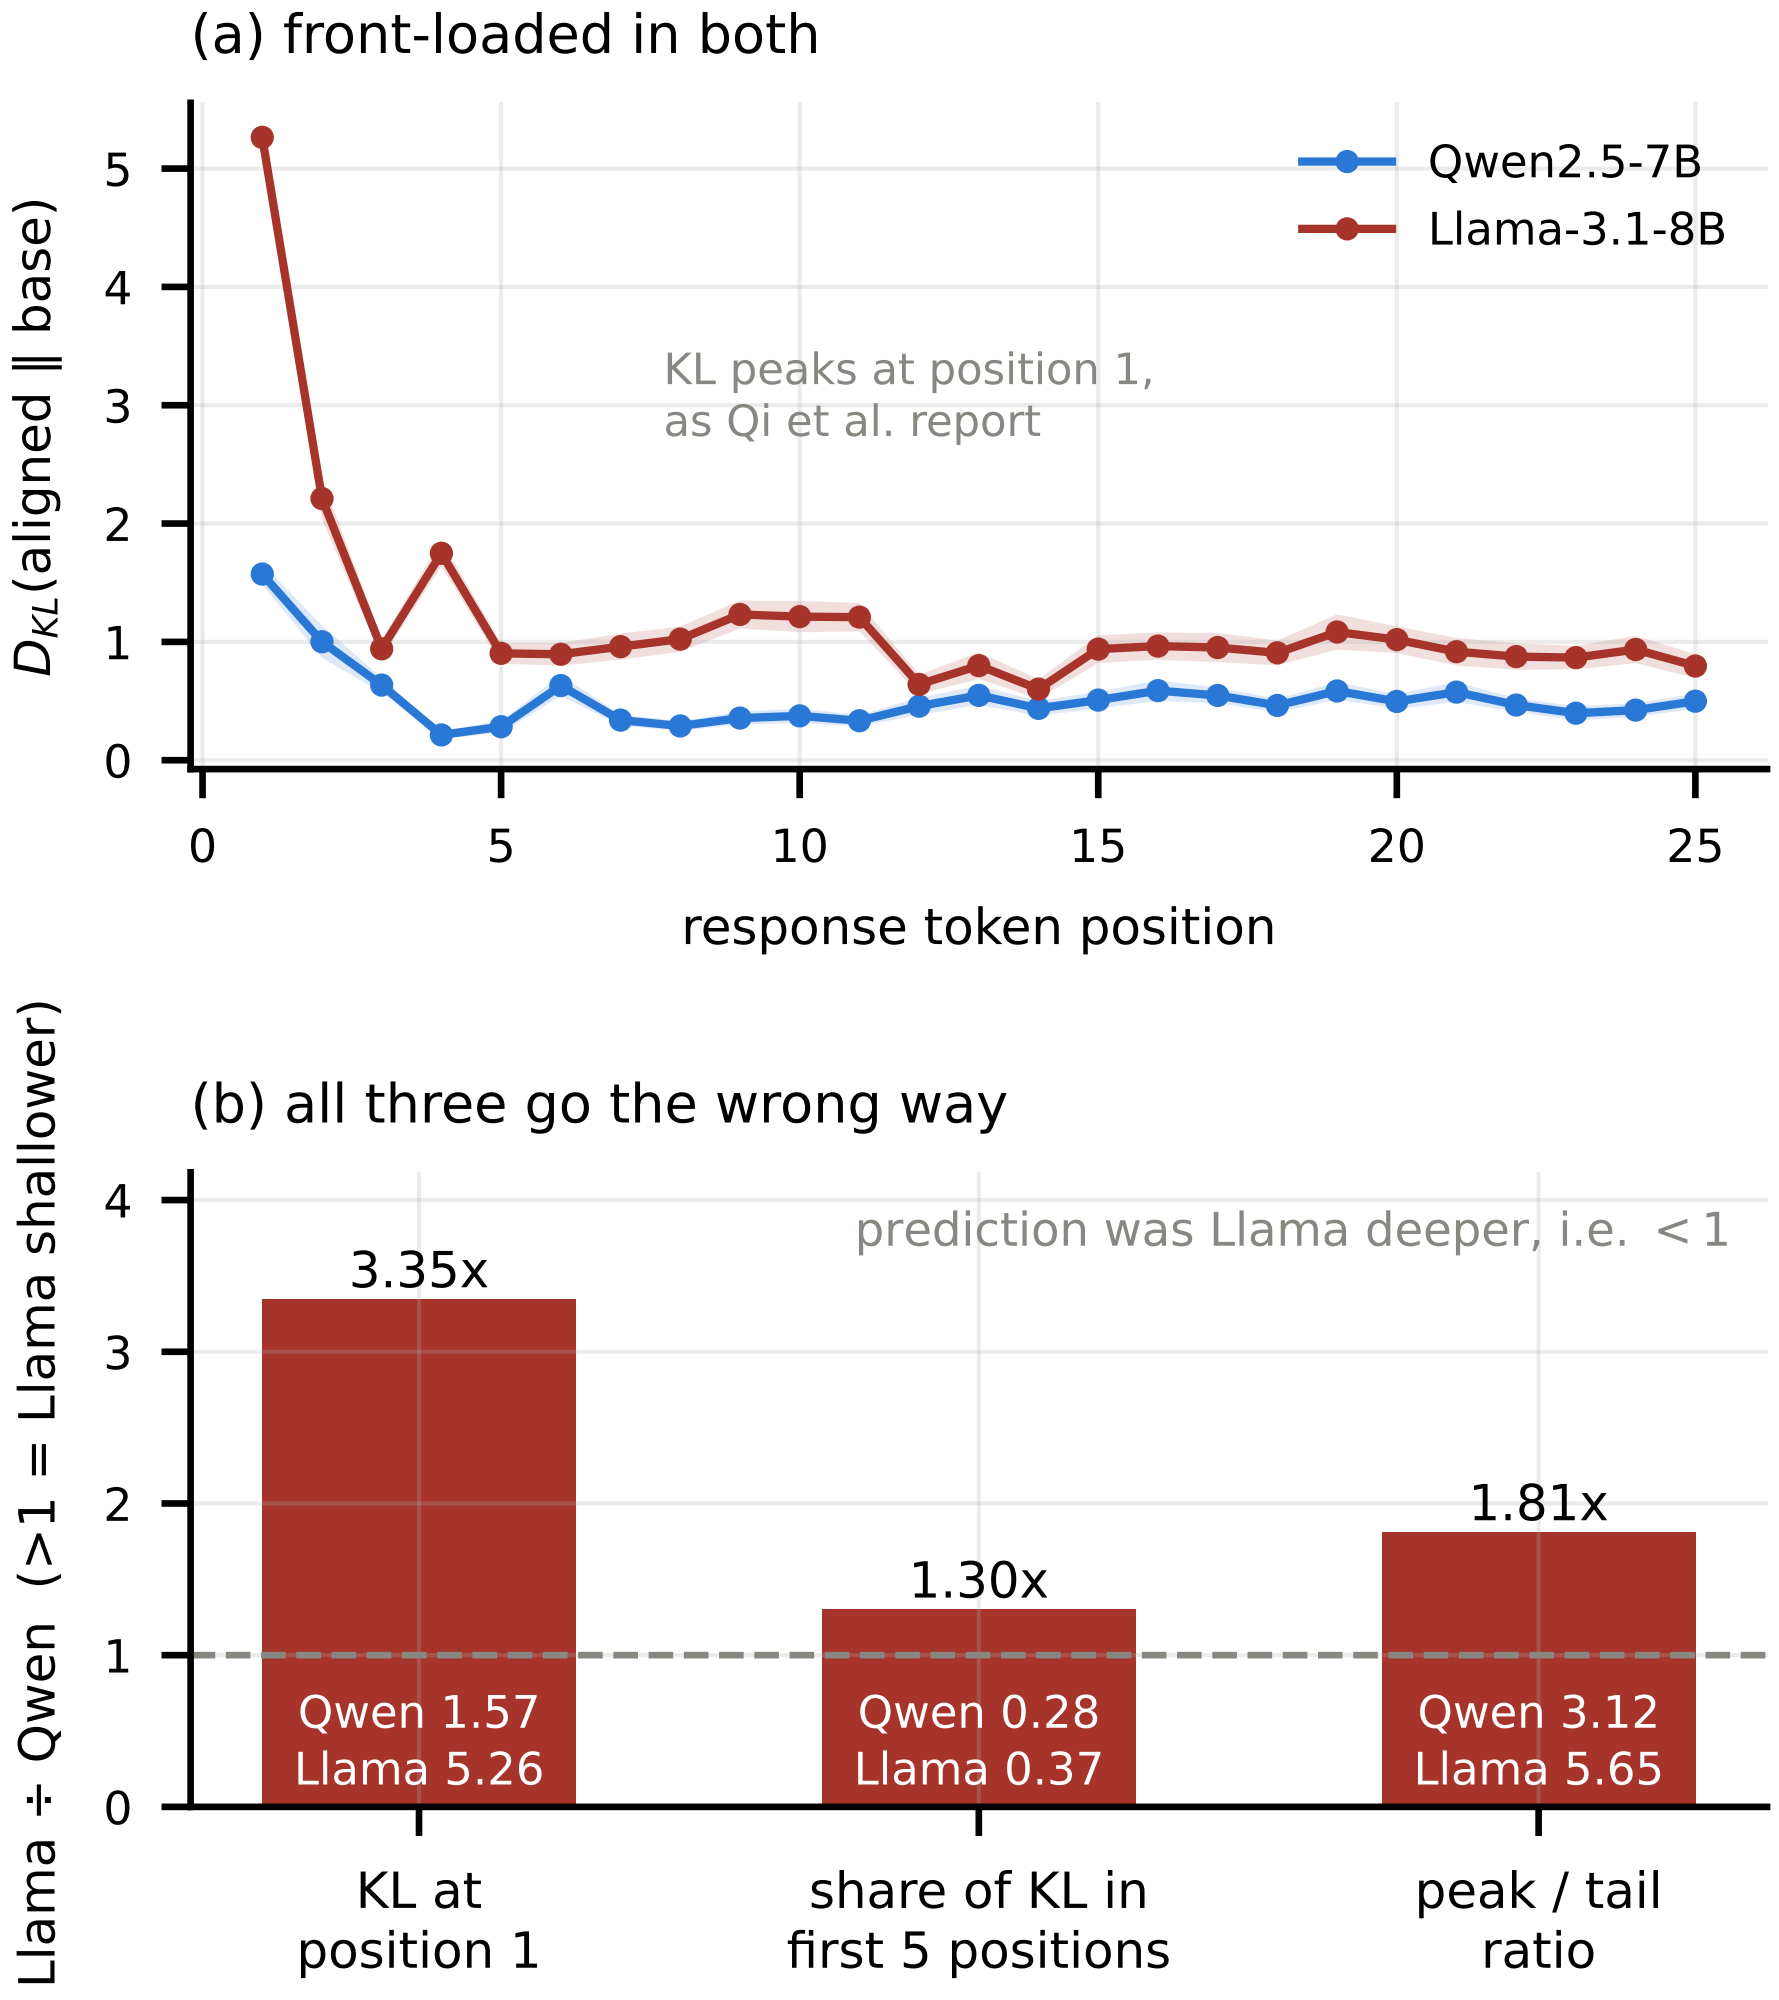
(a) front-loaded in both
D_{KL}(aligned ‖ base)
Qwen2.5-7B
Llama-3.1-8B
KL peaks at position 1,
as Qi et al. report
response token position
(b) all three go the wrong way
Llama ÷ Qwen (>1 = Llama shallower)
prediction was Llama deeper, i.e. < 1
3.35x
1.30x
1.81x
Qwen 1.57
Llama 5.26
Qwen 0.28
Llama 0.37
Qwen 3.12
Llama 5.65
KL at
position 1
share of KL in
first 5 positions
peak / tail
ratio